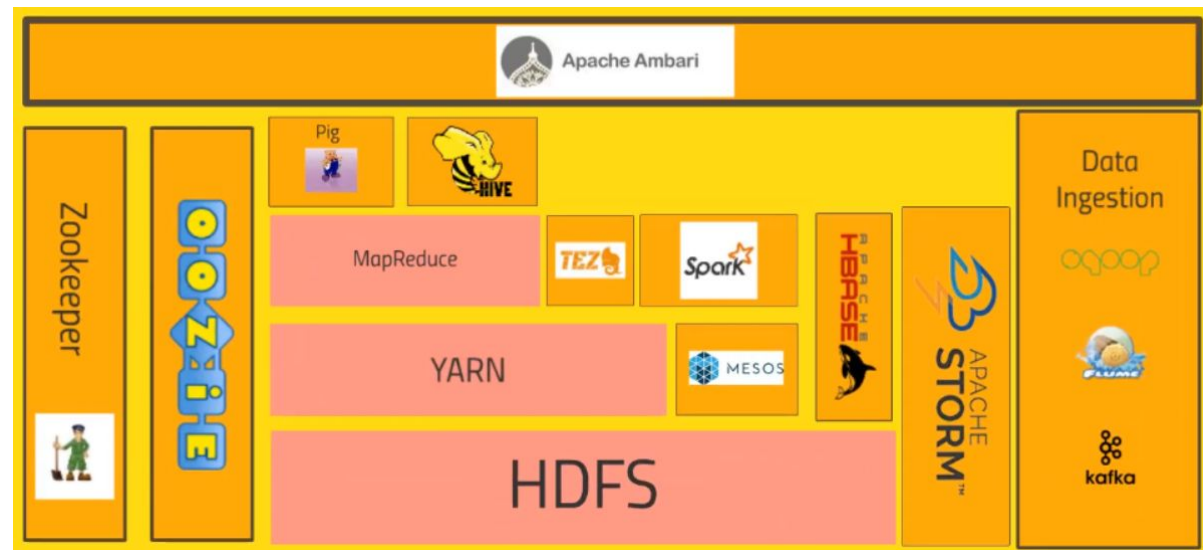


Hadoop & Distributed System Design

V2020.11.9
(Dr Yan Xu)

Hadoop Ecosystem



source: <https://www.udemy.com/course/the-ultimate-hands-on-hadoop-tame-your-big-data/>

1: Hadoop Core

- HDFS: a distributed file system
- YARN: a system to schedule applications
- MapReduce: a distributed algorithm framework

2: MapReduce Enhancement

- Pig: Scripts (pig latin) -> MapReduce
- Hive: SQL (hiveql) -> MapReduce
- Tez: a faster alternative of MapReduce

3: Spark

- In-memory processing
- Spark SQL + Spark Streaming

4: Databases

- HBase, Cassandra, MongoDB

5: Data Feeding

- Kafka, Flume, Sqoop

6: Analysing Data Streams

- Spark Streaming, Storm, Flink

7: Cluster Management

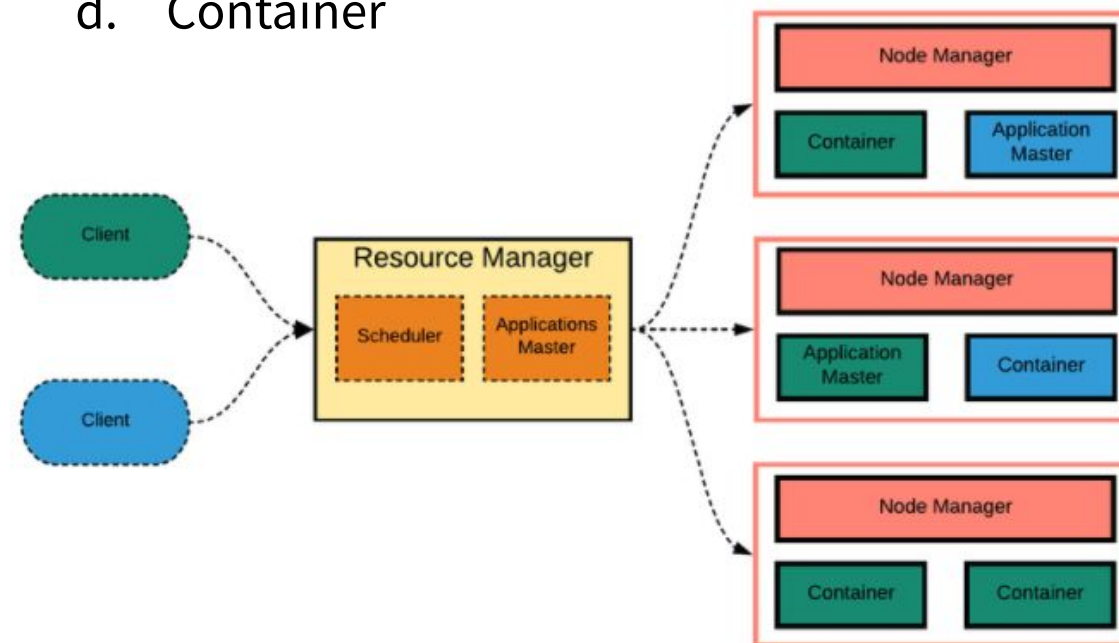
- ZooKeeper, Mesos, Oozie
- Zeppelin/Hue, Ambari

HDFS

- A master/slave architecture
 - NameNode** (master): manages the file system namespace and regulates access to files by clients
 - DataNodes** (slaves): manage storage attached to the nodes
- HDFS is built using the Java language
- Data Storage:
 - Block size
 - Block replication factor: e.g. 3
 - Rack-aware replica placement policy
- Transaction log (Metadata Persistence)
- Redundant NameNodes
- FS Shell
 - bin/hadoop dfs -mkdir /foodir
 - bin/hadoop dfs -cat /foodir/myfile.txt

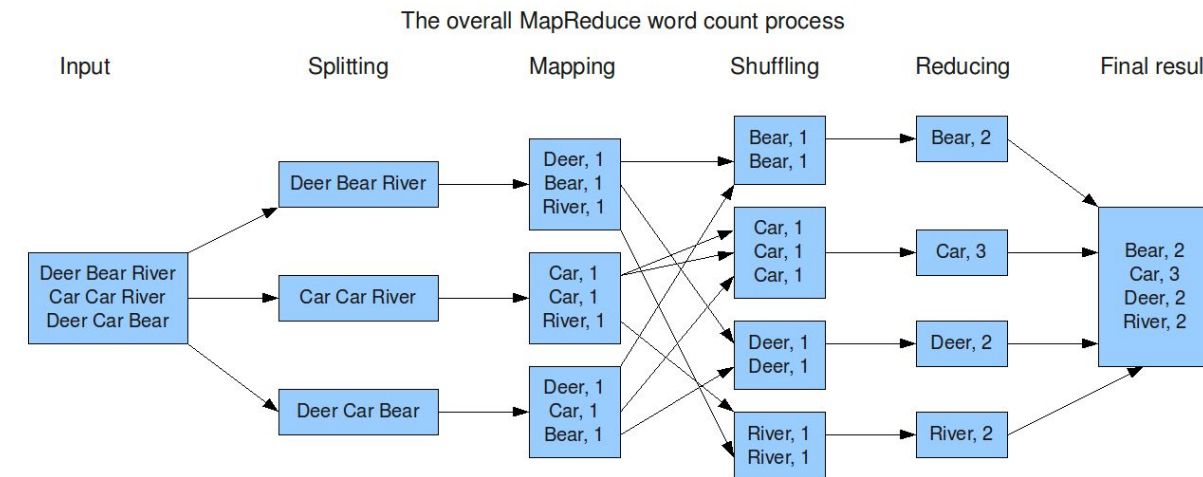
YARN

- Separates resource management layer from the processing layer, so to allows different data processing engines like graph processing, interactive processing, stream processing as well as batch processing on HDFS cluster.
- Main components:
 - a. Resource Manager
 - b. Node Manager
 - c. Application Master
 - d. Container



MapReduce

- A programming model to process big data on a cluster.
- Wordcount example:



- A MapReduce framework:
 - map()
 - shuffle(): redistribute data based on the output keys by the map func
 - reduce()
- MapReduce is designed to recover from the loss of whole nodes during the computation, so it writes interim results to distributed storage, which may not be a concern for your app and thus becomes not effective.


Hive

- **HiveQL**: gives an SQL-like interface to query data stored in various databases and file systems that integrate with Hadoop.
- In Hive, HiveQL is translated to a directed acyclic graph of MapReduce, Tez, or Spark jobs, which are submitted to Hadoop for execution.
- example:

```
1 DROP TABLE IF EXISTS docs;
2 CREATE TABLE docs (line STRING);
3 LOAD DATA INPATH 'input_file' OVERWRITE INTO TABLE docs;
4 CREATE TABLE word_counts AS
5 SELECT word, count(1) AS count FROM
6 (SELECT explode(split(line, '\s')) AS word FROM docs) temp
7 GROUP BY word
8 ORDER BY word;
```

Spark


- An open-source general-purpose distributed cluster-computing framework.



Details: <https://github.com/xuyannus/CheatSheet/blob/main/PySpark%20Basic%20Cheat%20Sheet.pdf>

Cassandra

- **Horizontal Scaling**: read and write throughput increase linearly as new machines are added
- **Masterless** (no single point of failure)
- Key Concepts:
 1. Gossip Protocol
 2. Consistent Hashing
 3. Wide-column Storage (or two-dimensional key-value store.)
 4. Cassandra Query Language (CQL)
 5. Shell: CQLSH



- Example CQL:

```
CREATE KEYSPACE MyKeySpace
WITH REPLICATION = { 'class' : 'SimpleStrategy', 'replication_factor' : 3 };

USE MyKeySpace;

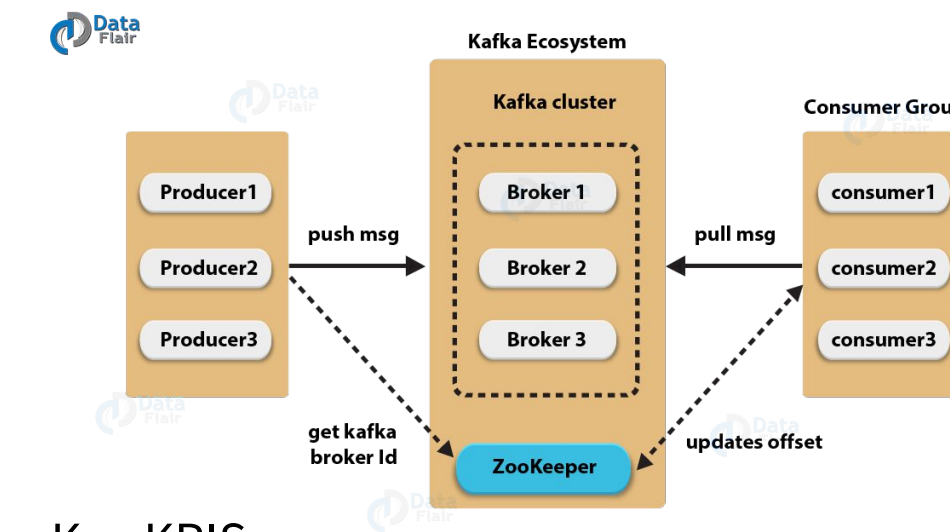
CREATE COLUMNFAMILY MyColumns (id text, Last text, First text, PRIMARY KEY(id));

INSERT INTO MyColumns (id, Last, First) VALUES ('1', 'Doe', 'John');

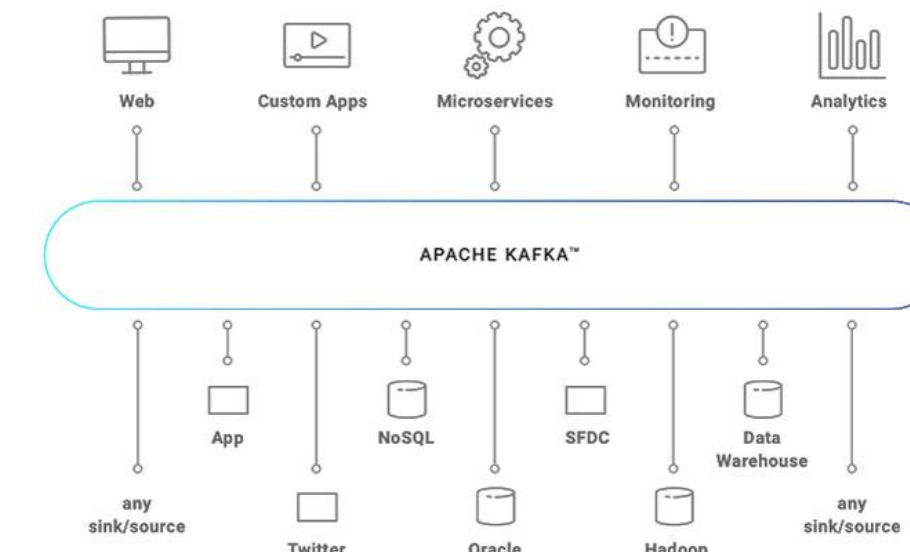
SELECT * FROM MyColumns;
```

Kafka

- to provide a unified, high-throughput, low-latency platform for handling real-time data feeds.
- Architecture



- Key KPIS:
 - Producer API
 - Consumer API
 - Streams API (input topics -> process -> output topics)
 - Connector API (connect to existing applications or data systems)
- Topic Partition
- Topic Replication Factor (e.g. 2)
- Application Scalability



Sqoop

- a command-line interface application for transferring data between RDBMS and Hadoop.
- Incremental load & Parallel load
- Commands:
 - sqoop import --connect jdbc:xx --table xx --target xx
 - sqoop export --connect jdbc:xx --table xx --export-dir xx
 - sqoop list
- Interactive shell: sqoop2-shell

Spark Streaming

- A streaming processing system that natively supports both batch and streaming workloads.
- **DStream**: RDDs with small batches
- **Microbatches** (e.g. 1 sec) & windowing (e.g. 30 mins)
- Example code:

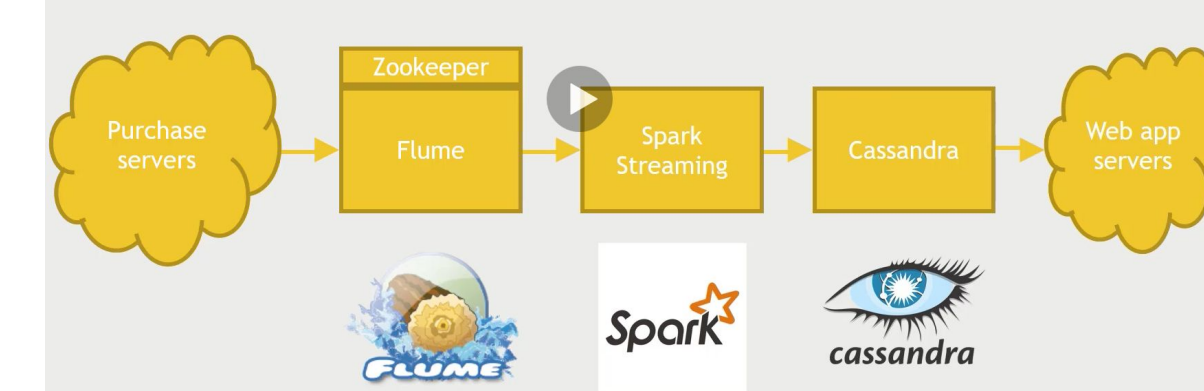
```
from pyspark import SparkContext
from pyspark.streaming import StreamingContext
from pyspark.streaming.kafka import KafkaUtils

sc = SparkContext(appName="PythonSparkStreamingKafka")
ssc = StreamingContext(sc, 5) # batch duration 5 sec
kafkaStream = KafkaUtils.createStream(ssc, kafka-stream-xxx) # kafkaStream is RDD
kafkaStream.map(xxx).countByValueAndWindow(60, 5) # window calculation
```

```
ssc.start()
ssc.awaitTermination(timeout=180)
```

System Design 1: top 10 best-selling books

Keywords: millions users, hourly update, low consistency

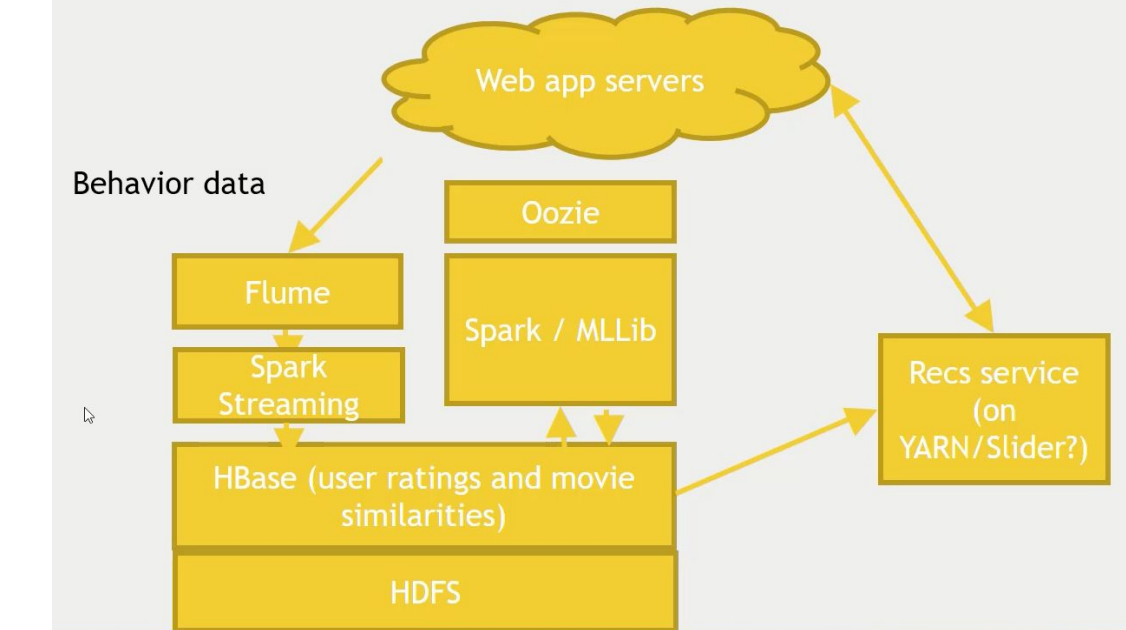


source: <https://www.udemy.com/course/the-ultimate-hands-on-hadoop-tame-your-big-data/>

1. Each Service itself (e.g. Spark, Cassandra) is running on a cluster

System Design 2: movie recommendation

Keywords: millions users, fast response to user behavior, low consistency



source: <https://www.udemy.com/course/the-ultimate-hands-on-hadoop-tame-your-big-data/>

1. ML model captures the more stable movie-movie similarity.
2. After a user action, Recs Service needs:
 - a. the user's recent actions (likes & dislikes)
 - b. movie-movie similarity to rank candidates and pick the top K movies