

Decision Tree Cheat Sheet

V2020.12.14
(Dr Yan Xu)

Motivation

1. Similar inputs have similar outputs
2. Similarity determined by a tree structure
 - could be considered as a faster and smarter K-NN

Goal

1. Minimize impurity (or loss) of leaves
2. Minimize # of leaves

Note: *Finding an optimal tree is NP-Hard*

Impurity (or Loss) measurement

1. Classification

- Gini impurity:

$$G(S) = \sum_{k=1}^c p_k (1 - p_k)$$

where, p_k is traction of inputs with label k

- Entropy (or, KL -Divergence to Uniform):

$$\sum_k p_k \log(p_k) + p_k \log(c)$$

where, c is number of classes

2. Regression

- Squared Loss:

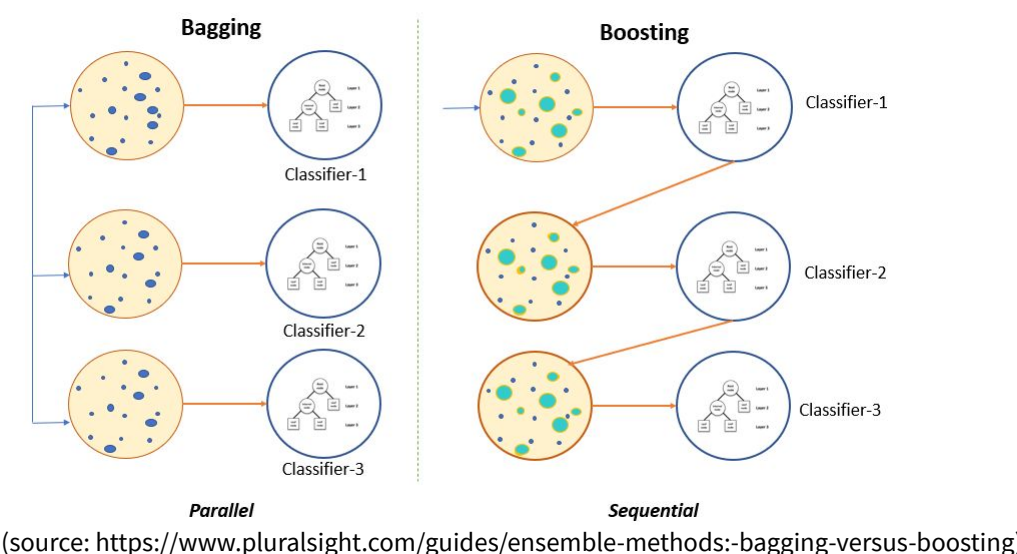
$$L(S) = \frac{1}{|S|} \sum_{(x,y) \in S} (y - \bar{y}_S)^2$$

Solution: Greedy

1. A top-down iterative split method
2. For each iteration, the split with the minimum impurity/loss is chosen
3. May stop early for some reasons (e.g. max depth)

Tree Ensembles

- One tree is not yet powerful, but ensembling trees is a great idea.
- Popular Solutions (Regression & Classification):
 - Random Forest
 - Gradient Boosting with Regularization (or XGBoost, LightGBM, etc)
- Bagging vs Boosting



Random Forest

- Bagging is a general method to deal with unstable predictions, so does Random Forest
- Original Algorithm [[paper](#)]:

1. Sample m data sets D_1, \dots, D_m from D with replacement.
2. For each D_j train a full decision tree $h_j()$ (max-depth= ∞) with one small modification: before each split randomly subsample $k \leq d$ features (without replacement) and only consider these for your split. (This further increases the variance of the trees.)
3. The final classifier is $h(\mathbf{x}) = \frac{1}{m} \sum_{j=1}^m h_j(\mathbf{x})$.

- Bootstrap Sample (with replacement) is used. Why? I am not sure.
- Subsampling features is optional.
- The number of trees is critical, but not sensitive, so typically easy to choose.
- Libraries: [scikit-learn](#)
- Variance Prediction: [forestci](#)
 - Be careful: it is sampling variance
- ML Code Example:
 - [RandomForest](#)
- Model Tuning: Easy

Gradient Boosting

- Residual Boosting
- For Square Loss, Gradient = Residual
- For other cost functions, Gradients are viewed as Residuals' generalization.
- Algorithm with Square Loss:

$$-g(x_i) = -\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} = y_i - F(x_i)$$

start with an initial model, say, $F(x) = \frac{\sum_{i=1}^n y_i}{n}$

iterate until converge:

calculate negative gradients $-g(x_i)$

fit a regression tree h to negative gradients $-g(x_i)$

$F := F + \rho h$, where $\rho = 1$

- Algorithm with General Cost Function

start with an initial model, say $F(x) = \frac{\sum_{i=1}^n y_i}{n}$

iterate until converge:

calculate negative gradients $-g(x_i) = -\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)}$

fit a regression tree h to negative gradients $-g(x_i)$

$F := F + \rho h$

GBM with Regularization (XGBoost)

- Regularized Cost Function:

$$\mathcal{L}(\phi) = \sum_i l(\hat{y}_i, y_i) + \sum_k \Omega(f_k)$$

where $\Omega(f) = \gamma T + \frac{1}{2} \lambda \|w\|^2$

- Second Order Approximation
 - faster convergence
- Shrinkage (or Step Size)
- Split Candidates: Quantile sketch
- Missing Value
 - default direction with higher gain
- Parallel & Distributed Computing
- Other Libraries:
 - [LightGBM](#), [CatBoost](#)
- ML Code Example:
 - [XGBoost](#)
- Model Tuning: Not Trivial