



PostgreSQL Cheat Sheet

(source: <https://www.udemy.com/course/sql-and-postgresql>)

V2021.03.14

(Dr Yan Xu)

Basic SQL

- SELECT:
 - **select** column1, column2, ...
 - **from** table_name
 - **where** condition;
- INSERT:
 - **insert into** table_name (column1, column2, ...)
 - **values** (value1, value2, ...);
- UPDATE:
 - **update** table_name
 - **set** column1 = value1, column2 = value2, ...
 - **where** condition;
- DELETE:
 - **delete from** table_name
 - **where** condition;

Intermedia SQL

- GROUP BY / ORDER BY
- LIMIT / OFFSET
- SET OPERATIONS:
 - **union & union all**
 - **intersect & intersect all**
- JOIN
 - **left join / right join**
 - **inner join (default) / outer join**
- IN, BETWEEN
- WINDOW Func
 - OVER (PARTITION BY xx ORDER BY xx DESC)
 - examples:
 - row_number(), rank()
 - Lag() - before the current row
 - Lead() - after the current row
- AS - table / feature rename

Scheme Design

- Entity Relationship Diagram
- Entities?
- Relationship?
 - 1-to-1:
 - 1-to-N / N-to-1: foreign key
 - N-to-N: a separate relationship table
- Primary Key (e.g. auto-generated ID or uuid)
- Foreign key & Deletion Protection (careful-to-use)
- Common Data Types:
 - integer, double precision, numeric
 - char, varchar, varchar(N), text
 - boolean, date, time, datetime, interval

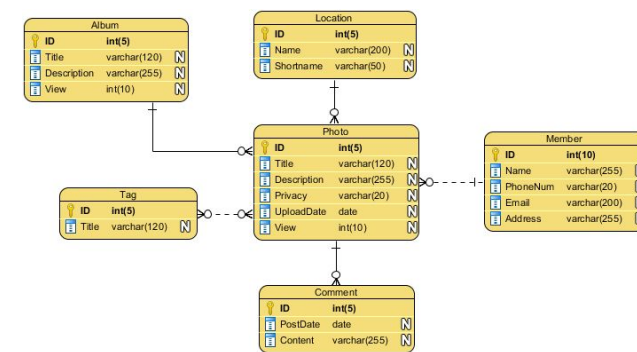
- Common attributes to consider:
 - created_at, updated_at, status

- Extra considerations:

- NOT NULL
- DEFAULT
- UNIQUE

- Criterion

- clear logic
- query performance
- less duplication



Scheme Management

- CREATE:
 - **create table** users (
 - id SERIAL PRIMARY KEY,
 - name VARCHAR NOT NULL,
 - price NUMERIC(10,2) DEFAULT 0);
 - **create index** idx_name ON users(name);
 - **create view** v(c1,c2) AS
SELECT c1, c2
FROM users;
- DROP:
 - **drop table** users
 - **drop index** idx_name
- ALTER:
 - **alter table** users **add** column
 - **alter table** users **add** constraint
 - **alter table** users **rename to** another_name

Database-side Verification & Constraints

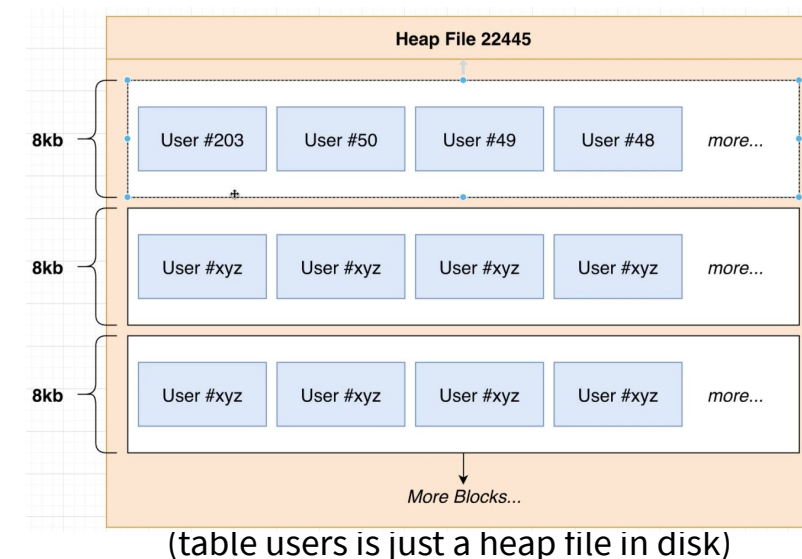
- DB-side vs App-side Verification
- Common Verification:
 - NOT NULL
 - DEFAULT
 - UNIQUE
 - MULTI-COLUMN UNIQUE
 - CHECK

- Example:

```
1 CREATE TABLE USERS (
2   id SERIAL PRIMARY KEY,
3   ssn VARCHAR(10) UNIQUE,
4   name VARCHAR(30) NOT NULL,
5   weight INTEGER CHECK(weight > 0)
6 );
```

PostgreSQL Internal

- Data Location
 - show data_directory;
- Data Storage
 - Heap file
 - File block
 - Records



- Index

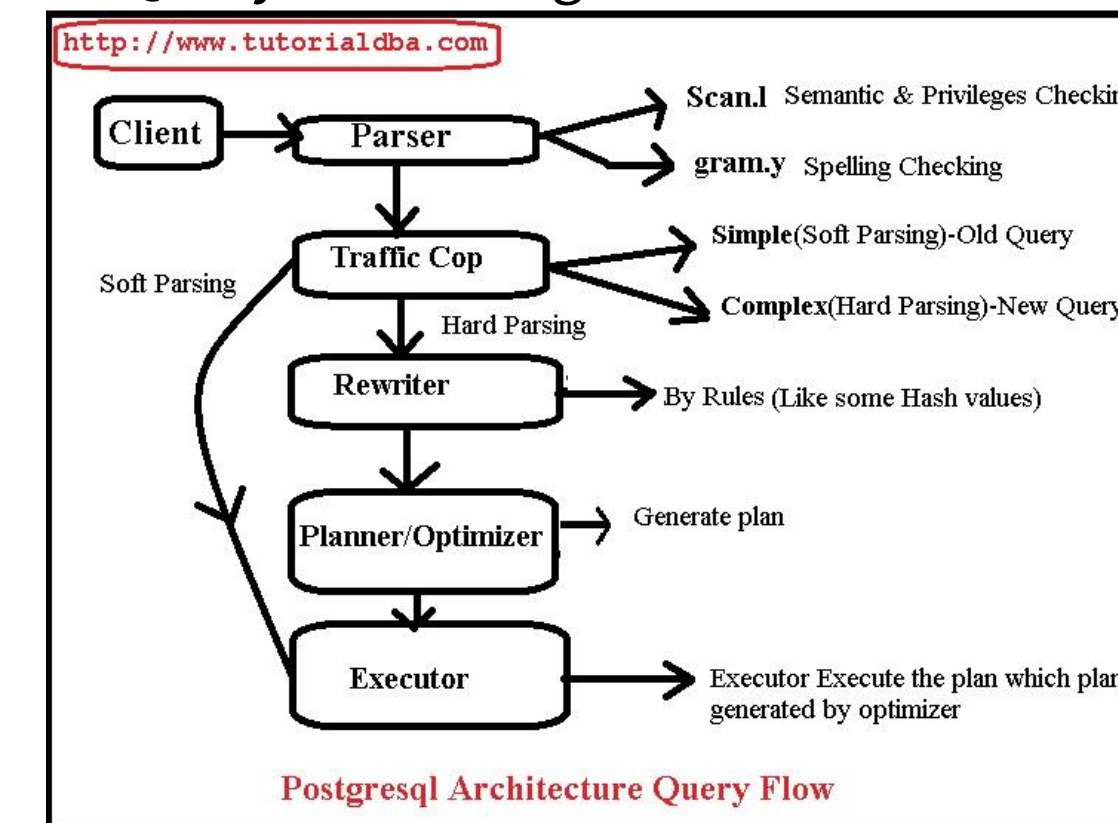
- Index is a carefully designed B+ Tree as a smaller table (and thus a smaller heap file)
- Postgresql automatically creates Index for Primary key and Unique columns.

- Metadata in DB

- pg_database
- pg_class: storage metadata
- pg_stat_activity: tracing & blame metadata
- pg_stat_*: used for query planning

Query Planning and Cost Analysis

- Query Processing Framework



- Commands:

- Explain
- Explain Analyze

- Query Plan



- Planner Cost Constants

- postgresql pre-defines the costs of most operations (e.g. an index read) relative to a sequential disc scan
- check details [here](#)

- Query Planning relies on pg_stats (it is a view) for Cost Estimation

- Command to update statistics:

- ANALYZE
- ANALYZE TABLE_NAME
- ANALYZE TABLE_NAME (COLUMN_NAME)

- Scanning large tables should be avoided if possible, by using indexes

Common Table Expression (with ...)

- **with** should be used to simplify SQL
- Format:

```
WITH
  engineers as (
    select *
    from employees
    where dept='Engineering'
  )
select *
from engineers
where ...
```

- **with** does not impact performance

Recursive Common Table Expression

- It should not be used
 - Problem: could easily be very costly

- Format:

```
with recursive ancestors as (
  select * from folks
  where name = 'Alex'
  union [all]
  select f.*
  from folks as f, ancestors AS a
  where f.id = a.father or f.id = a.mother
)
select * from ancestors;
```

Transaction

- All or Nothing

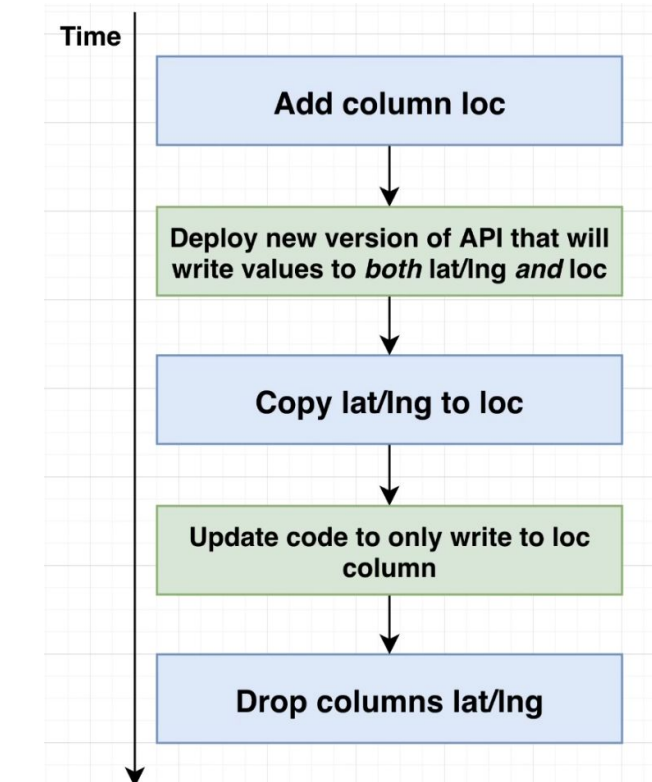
- The impact is within a local env before COMMIT
- It is similar with local branch in Git

- Format:

```
One transaction {
  BEGIN;
  UPDATE accounts SET balance = balance - 100.00
  WHERE name = 'Alice';
  UPDATE branches SET balance = balance - 100.00
  WHERE name = (SELECT branch_name FROM accounts
  WHERE name = 'Alice');
  UPDATE accounts SET balance = balance + 100.00
  WHERE name = 'Bob';
  UPDATE branches SET balance = balance + 100.00
  WHERE name = (SELECT branch_name FROM accounts
  WHERE name = 'Bob');
  COMMIT;
```

Schema / Data Migration Plan

- Any schema/data migration needs to have a PLAN, no matter how small the change is (e.g. column name change).
- Schema/Data migration may take days
- Migration Plan Example

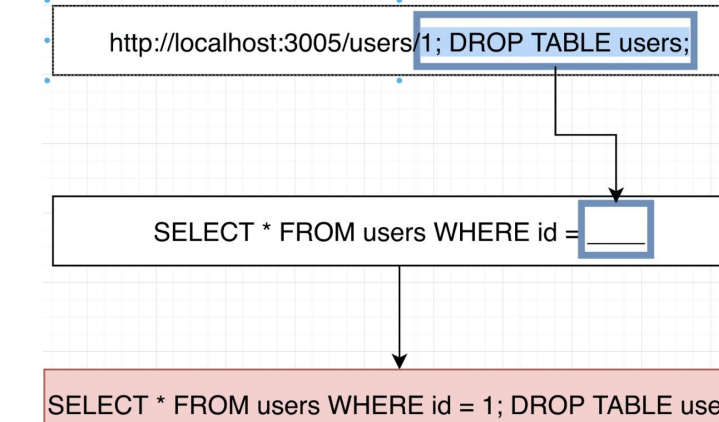


- Schema Migration as Code (Optional)

- python: [voyager-migrations](#)

Security

- SQL Injection Risk



- Prepared Statement (enforce single statement)

